

Machine Learning

# Building recommendation systems in Python with LightFM

Petrus Janse van Rensburg

Machine Learning

# Using LightFM to run a dairy farm in Tanzania

Petrus Janse van Rensburg

## Machine Learning

Using LightFM to ~~run a dairy~~  
~~farm in Tanzania~~ recommend  
groceries in Cape Town

Petrus Janse van Rensburg

**NETFLIX**



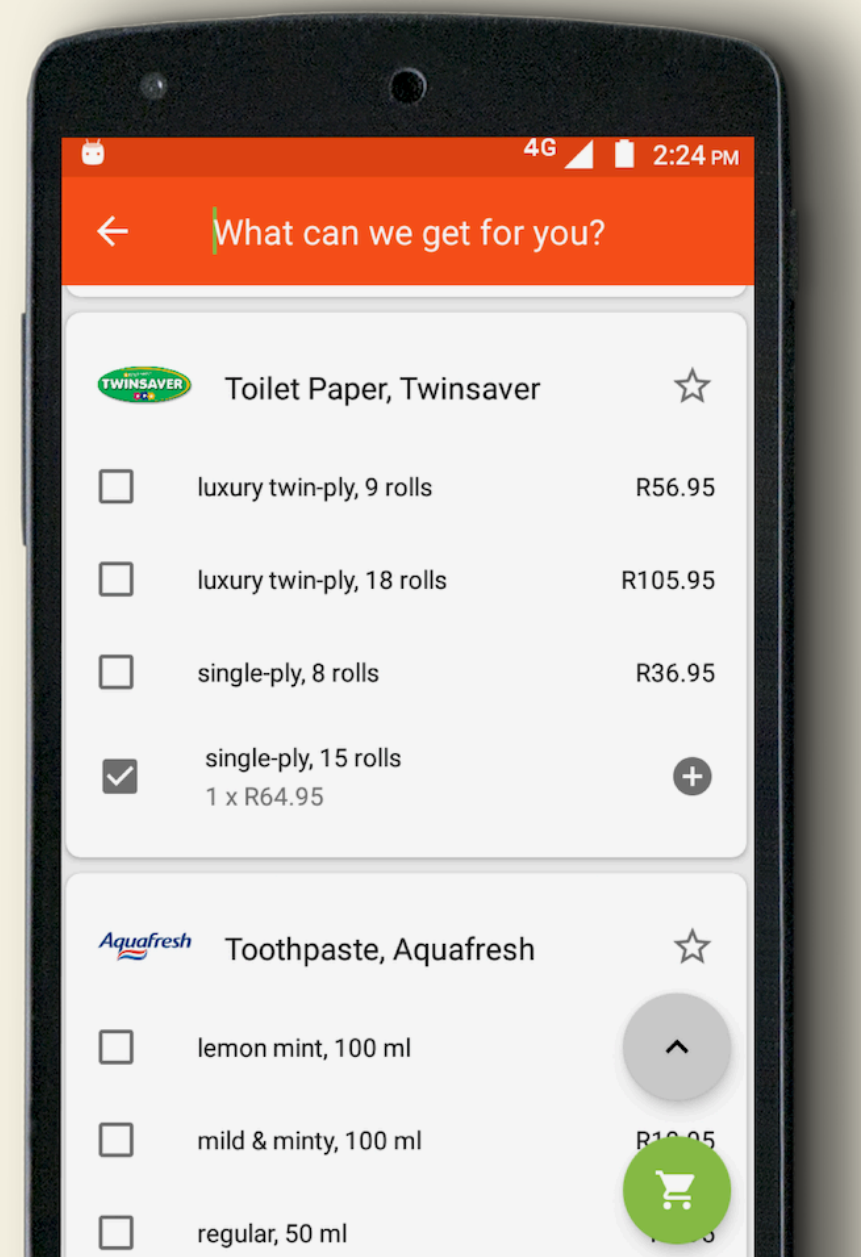
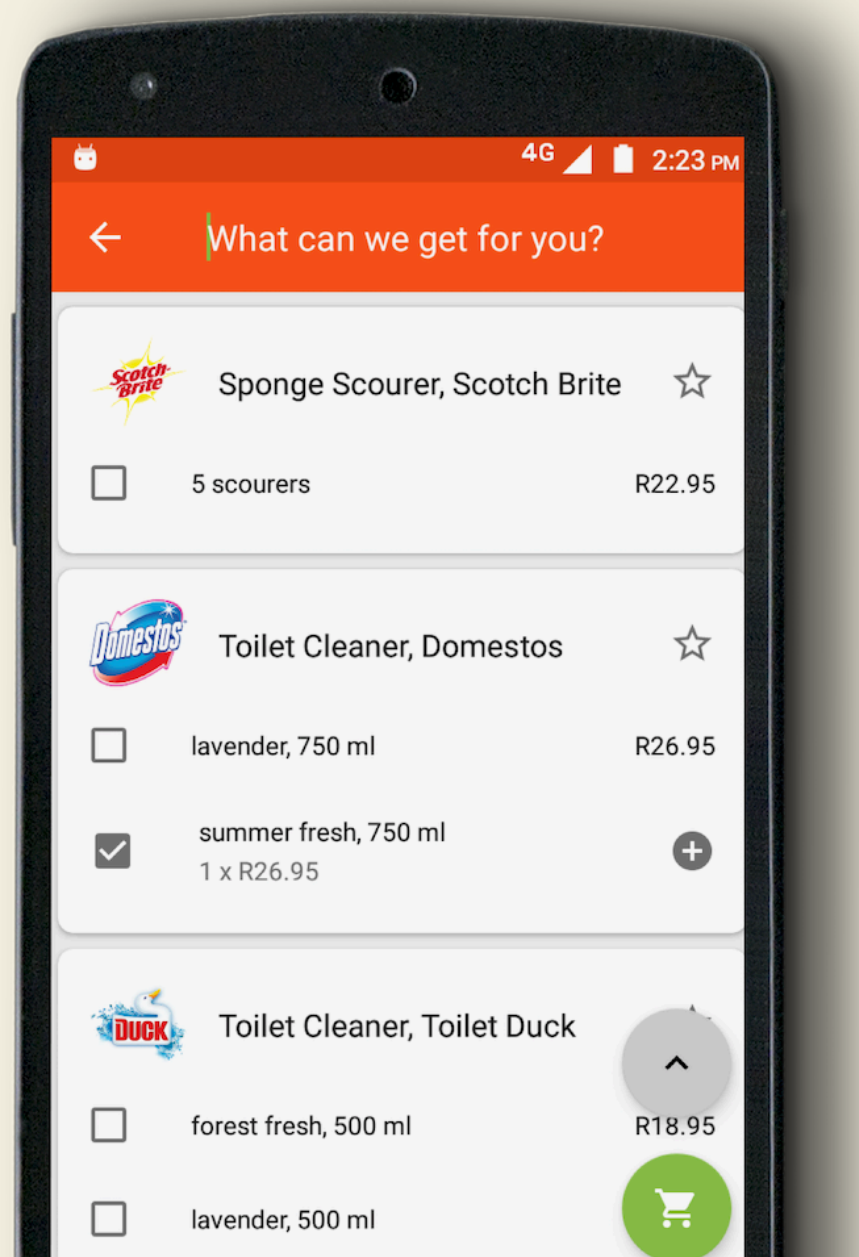
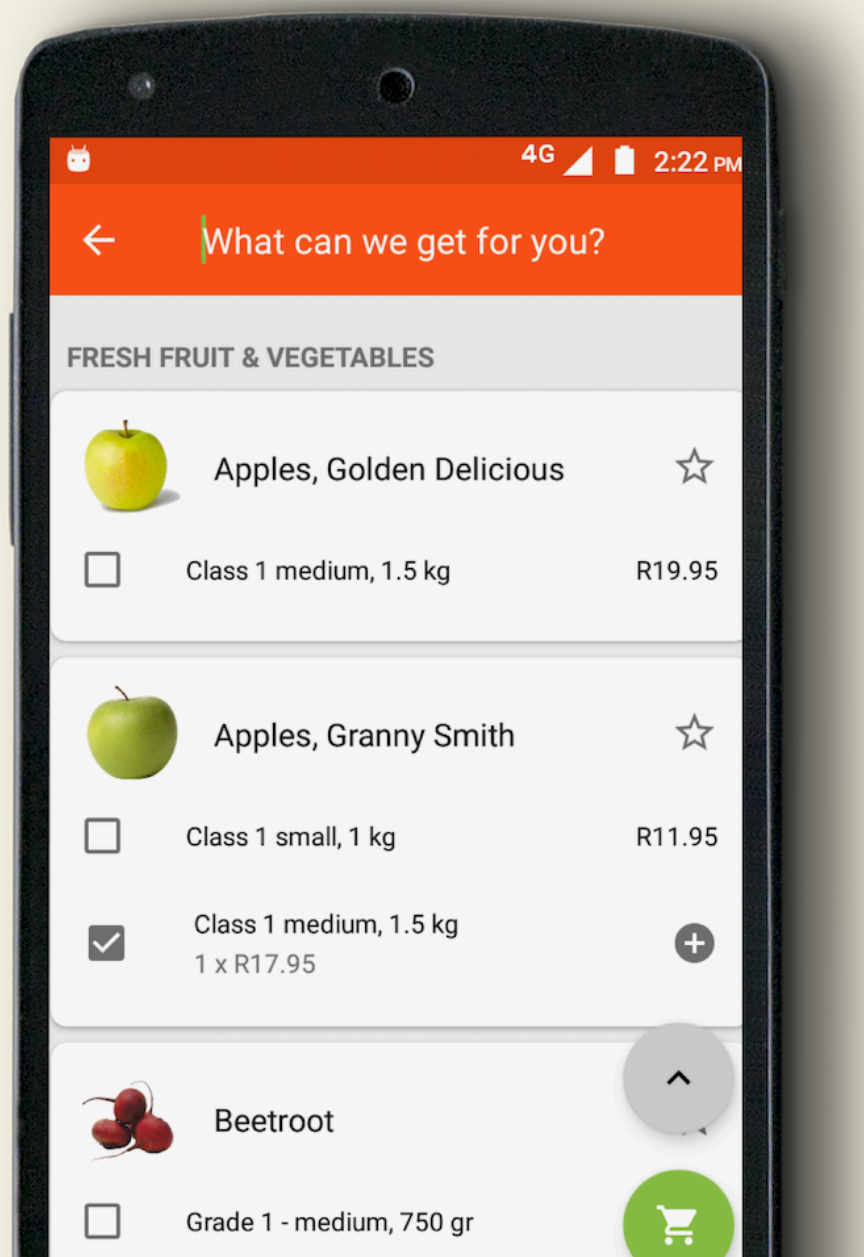
**UBER**  
**eats**

**amazon**

“People who liked this item, also liked...”

# LightFM

- Incorporates user-item interaction data + feature information into the same model
- E.g. collaborative + content-based filtering
- Uses a traditional matrix-factorisation approach (rather than a neural-network approach)
- Useful for a wide range of problems (movies, restaurants, groceries)



“It’s tough to make predictions, especially about the future”

– *Neils Bohr and others*



# Collaborative filtering

User-item interaction matrix

$$I = \begin{bmatrix} 1 & 0 & 0 & \dots & 0 & 0 & 0 \\ 0 & 0 & 0 & \dots & 1 & 0 & 1 \\ 0 & 1 & 0 & \dots & 1 & 0 & 0 \end{bmatrix}$$

Calculate cosine distances to find similar users and make recommendations based on their interactions

# Collaborative filtering

	five roses tea	huggies diapers	iwisa maize meal		knorrox stock	oros squash	tomatoes	
$I =$	1	0	0	...	0	0	0	user 1
	0	0	0	...	1	0	1	user 2
	0	1	0	...	1	0	0	user 3

Calculate cosine distances to find similar users and make recommendations based on their interactions

# Matrix factorization

Low-dimensional user- and item embeddings

$$\begin{matrix} I \\ (users \times items) \end{matrix} \approx \begin{matrix} U \\ (users \times n) \end{matrix} \begin{matrix} P \\ (n \times items) \end{matrix}$$

simple multiplication to predict scores

# Matrix factorization

Dot-product between user and item representations

$$\textit{score } 1 = \underset{\textit{(user } 1\textit{)}}{[0.25 \quad 0.93 \quad 1.43 \quad 0.48 \quad 0.34]} \underset{\textit{(item } 1\textit{)}}{\begin{bmatrix} 0.53 \\ 0.34 \\ 0.53 \\ 0.82 \\ 1.37 \end{bmatrix}}$$

# Matrix factorization

- You can decide how many dimensions to use for your embeddings.
- Factorization can be done in a number of ways, e.g.
  - Stochastic Gradient Descent
  - Alternating Least Squares

# Shortcomings

- Some users haven't bought anything.
- Some products haven't been bought.

# Content-based filtering

Make use of item feature information

*['huggies', 'diapers']*

*['huggies', 'nappy pants']*

*['huggies', 'baby wipes']*

*['pampers', 'diapers']*

One-hot encode item features, then calculate cosine distances to find similar items and make recommendations

# One-hot encoding

Vectorize categorical variables

	huggies	pampers	diapers	nappy pants	baby wipes
<i>item 1</i>	1	0	1	0	0
<i>item 2</i>	1	0	0	1	0
<i>item 3</i>	1	0	0	0	1
<i>item 4</i>	0	1	1	0	0



# Shortcomings

- Doesn't work for new users.
- No information-sharing between users.

# Hybrid approach

user-item interactions

+

item features

+

user features

# Hybrid approach

Users can also have features

*['Cape Town', 'Observatory']*

*['Cape Town', 'Mitchell's plain']*

*['Cape Town', 'Langa']*

*['Cape Town', 'City Center']*

# One-hot encoding

Vectorize categorical variables

cape town  
observatory  
mitchell's plain  
langa  
city center

$$user\ 1 = [1\ 1\ 0\ 0\ 0]$$

$$user\ 2 = [1\ 0\ 1\ 0\ 0]$$

$$user\ 3 = [1\ 0\ 0\ 1\ 0]$$

$$user\ 4 = [1\ 0\ 0\ 0\ 1]$$

# LightFM

## How it works?

- Matrix Factorization that can incorporate interactions, item features and user features.
- Feature representations are summed before taking a dot-product for calculating score.
- Reduces to regular MF model when no features are provided.

# LightFM

## Matrix Factorization with Stochastic Gradient Descent

- Loss Function

Bayesian Personalised Ranking (BPR)

Weighted Approximate-Rank Pairwise (WARP)

- Learning Rate

- No. of components in latent space

# LightFM

$$\begin{matrix} I \\ (user\ features \times item\ features) \end{matrix} \approx \begin{matrix} U \\ (user\ features \times n) \end{matrix} \begin{matrix} P \\ (n \times item\ features) \end{matrix}$$

addition, then multiplication to predict scores

# LightFM

Addition, then multiplication to predict scores

*user 1 = western cape + cape town + observatory*

*item 1 = huggies + diapers*

*score = user 1 · item 1*



# Testing / evaluation

Split your interactions dataset, keep one part out for testing.  
The goal is to not have a large discrepancy between train and test sets (otherwise you're overfitting).

1. AUC ROC score:  
(Area-Under-Curve Receiver Operating Characteristic)  
The probability that a randomly chosen positive example has a higher score than a randomly chosen negative example. A perfect score is 1.0
2. Precision at k score:  
The fraction of known positives in the first k positions of the ranked list of results. A perfect score is 1.0.

# Grocery recommendation example

# Assumptions

1. The future will be like the past.
2. Synonymy is accounted for when extracting features e.g.
  - “bath soap” vs. “beauty soap”
  - “peanut butter” vs. “butter”
3. Time-sensitivity is accounted for.

# Further reading

- LightFM 'examples' directory on GitHub
- Maciej Kula - PyData 2016 talk
- Maciej Kula article on arxiv: "Metadata Embeddings for User and Item Cold-start Recommendations"
- Spotlight package, if you'd rather use neural network / deep-learning approach

# Thank you

[petrus.jvrensborg@gmail.com](mailto:petrus.jvrensborg@gmail.com)  
[@petrusjvr](#)